



PRODUCT ROADMAP · INVESTOR BRIEF

Pre-Market Enhancement Roadmap

Nine enhancement pillars to convert RegGuard AI v2.2 from a working technical demonstration into an enterprise-grade, investable compliance automation platform — sequenced for impact ahead of formal market assessment.

PILLARS

09

HORIZON

4Q

VIEWS LIVE

29

FEEDS

04

Table of Contents

01	Executive Summary	03
02	Current State Snapshot (v2.2 Baseline)	04
03	Strategic Thesis & Market Window	05
04	Pillar 1 — AI Intelligence Layer	06
05	Pillar 2 — Workflow & Collaboration	08
06	Pillar 3 — Evidence Vault & Audit Readiness	09
07	Pillar 4 — Broader Data Coverage	10
08	Pillar 5 — Risk & Screening Automation	11
09	Pillar 6 — Integration & Public API	12
10	Pillar 7 — Security & Trust (Enterprise Sales Unlock)	13
11	Pillar 8 — UX Polish	14
12	Pillar 9 — AI Governance & Explainability	15
13	Impact vs Effort Prioritization Matrix	16
14	Quarterly Delivery Timeline	17
15	Top 5 Pre-Market Recommendations	18

— 01 • EXECUTIVE SUMMARY

From Technical Demonstration to Investable Platform

RegGuard AI v2.2 has reached a meaningful inflection point. The platform ships **29 state-machine views** across six functional zones, four of which are wired to real free-tier regulatory data feeds (Federal Register, ESMA RSS, OFAC SDN, EU CFSP, GLEIF LEI). Audit entries are anchored to a real

SHA-256 Merkle root with optional Polygon Amoy testnet broadcast, and the Network Graph Explorer runs a custom force-directed simulation with full mobile touch support. The “machine proposes, human confirms” design principle is enforced consistently across every view via the shared **BooleanActionCard** component.

However, technical capability alone does not translate into enterprise revenue. Before formal market assessment can begin, three conversion gaps must be closed. First, the AI layer is currently informational rather than operational — it surfaces data but does not generate grounded recommendations or automate policy gap analysis. Second, the platform lacks the enterprise security primitives (SSO, SOC 2, audit packaging) that procurement teams require before signing contracts above \$50K ARR. Third, the data layer is US/OFAC-centric, which limits the addressable market to roughly one-third of the global GRC opportunity.

This document proposes **nine enhancement pillars** organized into three investment tiers. Tier 1 (Pillars 1, 3, 7) addresses the highest-leverage gaps: AI moat depth, audit-ready evidence vault, and SOC 2 / SSO for enterprise sales unlock. Tier 2 (Pillars 4, 5, 6) expands addressable market through jurisdiction coverage, screening depth, and a public REST API enabling platform partnerships. Tier 3 (Pillars 2, 8, 9) delivers workflow, UX, and AI governance polish that compound the value of Tier 1 and Tier 2 investments.

TARGET OUTCOME

By end of H2 next year, RegGuard AI should achieve a **4x improvement** in three priority dimensions: AI moat depth (informational → grounded & actionable), enterprise readiness (no SSO → SOC 2 Type II + SSO + audit packages), and platform extensibility (read-only dashboard → public REST API with usage metering). These three shifts convert the platform from a technical showcase into a credible Series A investment opportunity.

The remainder of this document walks through each pillar in detail, including technical approach, effort estimate in engineer-weeks, and the strategic rationale. The closing two sections provide a visual impact-vs-effort prioritization matrix and a quarter-by-quarter delivery timeline covering Q3 this year through Q2 next year.

02 • BASELINE

Current State Snapshot (v2.2)

Before proposing enhancements, it is essential to anchor on what RegGuard AI v2.2 has already shipped. The table below quantifies the baseline across the six dimensions that matter most to enterprise buyers and investors. The numbers reflect the live deployment at the project’s GitHub Pages URL, not a roadmap aspiration.

VIEWS LIVE

29

across 6 zones

REAL FEEDS

04

Federal Register, ESMA, OFAC, EU CFSP, GLEIF

PRISMA MODELS

29


SQLite, streaming-ready

CHAIN ANCHORS

16

Merkle root + Polygon Amoy

The capability radar below scores RegGuard AI on six axes, with each axis rated on a 1-to-5 maturity scale where 5 represents enterprise-grade. The amber polygon is current state; the dashed cyan outline is the target after the enhancements in this roadmap are delivered. Audit capability is the strongest axis today (level 4) thanks to the existing Merkle anchoring work. Integration is the weakest (level 1) because no public API or third-party connectors exist yet.



Three observations frame the rest of this document. First, the **data layer is mature for US/OFAC** but shallow for other jurisdictions — this caps the addressable market. Second, the **AI layer is informational** (it surfaces regulations and entities) but does not yet generate grounded recommendations, gap analyses, or drafted policy language. Third, the **security and integration layers are at level 1-2**, which is the single biggest blocker to closing enterprise deals above pilot pricing.

Market Window & Competitive Position

The global Governance, Risk & Compliance (GRC) software market is estimated at **\$36B in 2025** and projected to reach \$64B by 2027, growing at roughly 12% CAGR. The AI-native compliance sub-segment is growing 3-4x faster than the overall market, driven by three converging forces that create a rare window for an AI-first entrant.

Force 1 — Regulatory velocity outpacing manual capacity. Federal Register publishes 3,000-4,000 final rules per year, ESMA issues 200+ opinions and guidelines annually, and global sanctions list updates occur daily. No compliance team can manually track this volume; AI-driven monitoring is becoming a survival requirement, not a productivity enhancement.

Force 2 — EU AI Act enforcement creates mandatory AI governance tooling demand. The EU AI Act's risk-tier framework, with enforcement beginning phased rollout this year and next, makes AI governance tooling a regulatory requirement rather than a nice-to-have. Financial services firms using AI for credit scoring, fraud detection, or customer segmentation must demonstrate conformity — creating direct demand for the governance capabilities proposed in Pillar 9.

Force 3 — Enterprise buyers shifting from GRC suites to composable AI-native platforms. Incumbent GRC suites (ServiceNow GRC, Archer, MetricStream) are workflow tools with thin AI layers bolted on. Modern buyers increasingly prefer composable platforms that combine native AI, real-time data feeds, and modern APIs — exactly the positioning RegGuard AI is built for.

DIFFERENTIATOR

RegGuard AI's "**machine proposes, human confirms**" design principle is the single most defensible positioning against both incumbent GRC suites (which lack native AI) and emerging AI-only compliance startups (which lack human accountability constructs). Regulators care about human accountability; operations teams care about AI scaling. This principle addresses both.

The nine enhancement pillars in the following sections are sequenced to capture this window. Tier 1 investments (Pillars 1, 3, 7) reinforce the AI moat and unlock enterprise sales. Tier 2 investments (Pillars 4, 5, 6) expand addressable market and enable platform partnerships. Tier 3 investments (Pillars 2, 8, 9) compound the value of Tier 1 and 2 by making the platform operationally useful, visually polished, and regulator-friendly.

— 04 · PILLAR 1

AI Intelligence Layer

The deepest moat — converts RegGuard from a data dashboard into a reasoning platform.

The AI Intelligence Layer is the single highest-impact investment in this roadmap. RegGuard AI today surfaces regulatory data and entity matches but does not reason over them. This pillar adds five

capabilities that together convert the platform from an informational tool into a grounded reasoning system — the kind of moat that incumbent GRC suites cannot replicate without rebuilding their data layer from scratch.

1.1 Regulatory Impact Analyzer

10 EW

Paste a new regulation (or point RegGuard at a Federal Register URL) and receive an auto-generated **gap analysis** against existing internal policies, scored by severity (critical / material / informational). Each gap includes a recommended remediation step and a traceable citation back to the regulation text. Technical approach: embedding similarity over a vector store of policy chunks, with an LLM-generated gap summary constrained to grounded citations.

1.2 Multi-Jurisdiction Diff Engine

8 EW

Compare the same regulatory concept (e.g. beneficial ownership threshold, AML transaction reporting limit) across **EU, US, UK, and APAC jurisdictions**, surfacing arbitrage opportunities and conflicting obligations. Particularly valuable for multinational financial institutions that must reconcile divergent regimes. Builds on the existing Regulation Watch infrastructure.

1.3 Semantic Regulation Search

6 EW

Embedding-based search across the full regulation corpus — not keyword matching. Users ask natural-language questions ("what are our obligations for crypto-asset custody?") and receive ranked passages with **cited source paragraphs** and a deep-link back to the original text. Underlying model: a sentence-transformer encoder with periodic re-indexing on regulation updates.

1.4 Compliance Q&A Copilot

9 EW

Upgrade the existing Assistant view from a generic chatbot into a **RAG-powered compliance assistant** that grounds every answer in cited regulations and internal policies. Critical design rule: the copilot must always show its sources and must refuse to answer when grounding is unavailable. This refusal behavior is itself a regulator-facing feature — it demonstrates that the AI is accountable, not hallucinating.

1.5 Auto-Generated Compliance Checklists

9 EW

Extract actionable obligations from regulation text and convert them into **assignable checklists** with owners, due dates, and evidence requirements. Closes the loop between regulation ingestion and operational execution. Builds on Pillar 2 (Workflow) for assignment and SLA tracking — the two pillars should be sequenced together.

REVENUE IMPLICATION

This pillar alone could justify a **2-3x ARR uplift** versus the current feature set, because it shifts RegGuard from a “read-only intelligence dashboard” (20 — 40K ARR per mid — market customer) into a “reasoning copilot” (60-120K ARR per enterprise customer). The gap-analyzer and checklist features map directly to budget lines that compliance teams already fund.

Total estimated effort for Pillar 1: **42 engineer-weeks** across the five features, plus a shared infra investment of approximately 6 engineer-weeks for the vector store, embedding pipeline, and citation framework. Recommended team: 2 ML/LLM engineers + 1 backend engineer + 1 frontend engineer, working in parallel across the five features. Realistic delivery: Q3 dev → Q4 Copilot v1 GA → Q2 next year full layer GA.

— 05 • PILLAR 2

Workflow & Collaboration

Converts RegGuard from a dashboard into an operational system of record.

Today, RegGuard AI surfaces insights but does not act on them. The “machine proposes, human confirms” pattern is implemented at the UI level via the BooleanActionCard component, but confirmed actions do not flow into assignments, SLAs, or audit trails. This pillar closes that loop, converting RegGuard from a tool that compliance teams look at into a tool they work inside of every day. Enterprise buyers consistently rank this capability as a top-three purchase criterion.

2.1 Approval Workflows with Multi-Stage Sign-Off

6 EW

Configurable approval chains (e.g. analyst → manager → CCO) with full audit trails. Each stage records the reviewer, timestamp, decision rationale, and a hash anchor extending the existing Merkle chain. Critical for regulator examinations: every decision is reconstructable end-to-end.

2.2 Task Assignment + SLA Tracking with Auto-Escalation

5 EW

Convert AI recommendations into assigned tasks with due dates and SLA policies. Overdue tasks **auto-escalate** to the next reviewer in the chain. This single feature transforms the daily operations of a compliance team — it eliminates the spreadsheet and email round-trips that today consume 30-40% of analyst time.

2.3 Comments & Annotations on Regulation Clauses

4 EW

Inline commenting anchored to specific regulation paragraphs, with @-mention notifications. Creates a persistent institutional knowledge layer — when a regulator asks “why did you interpret this clause this way?”, the answer is one click away.

2.4 Notification Channels (Email / Slack / Teams / Webhook)

4 EW

Configurable per-user and per-event-type notifications. Webhook channel enables downstream automation (e.g. auto-create a Jira ticket when a high-severity regulation change is detected).

2.5 Calendar Integration for Filing Deadlines (iCal Export)

3 EW

Every regulation with a filing deadline automatically appears on the user’s Outlook / Google Calendar via an iCal feed. Eliminates the “missed deadline” failure mode that drives the most expensive compliance incidents.

Total estimated effort for Pillar 2: **22 engineer-weeks**. Recommended team: 1 backend engineer + 1 frontend engineer, sequenced after Pillar 1 because the workflow layer depends on AI-generated recommendations to be useful. Realistic delivery: Q3 dev → Q4 GA.

— 06 · PILLAR 3

Evidence Vault & Audit Readiness

The single feature every compliance team buys — direct revenue driver, low technical risk.

The Evidence Vault extends RegGuard’s existing Merkle anchoring work into a complete audit-readiness system. Today, the chain evidence view anchors audit log entries to a Merkle root on Polygon Amoy. This pillar adds the surrounding capabilities that turn anchored logs into **regulator-ready audit packages** — the artifact every compliance team must produce during an examination. Most of the infrastructure already exists; the work is primarily UX and packaging.

3.1 Centralized Evidence Storage with Hash-Anchored Timestamps

4 EW

A document store for proof-of-compliance artifacts (policy documents, training records, attestation screenshots, configuration exports). Every artifact is automatically hashed and anchored to the existing Merkle chain, providing cryptographic proof that the artifact existed at a specific time and has not been modified since.

3.2 One-Click Regulator-Ready Audit Packages

5 EW

Generate a structured PDF / Excel bundle grouping evidence by examination scope (e.g. "AML program review", "sanctions screening effectiveness"). Each artifact in the bundle carries its Merkle proof and a verification URL. This is the single feature compliance teams reference most when explaining why they bought the platform.

3.3 Historical Trend Analysis of Compliance Posture

3 EW

Show how key compliance metrics (open risk items, policy review currency, training completion rates) evolved over the past 6-18 months. During an examination, this proves continuous improvement — a story regulators reward.

3.4 Heat Maps by Jurisdiction / Business Unit / Risk Category

2 EW

Visual heat maps identifying where compliance posture is weakest across the organization. Drives prioritization conversations between the CCO and business unit leaders — positioning RegGuard as the system of truth for the entire compliance program.

WHY THIS IS THE EASIEST SELL

Every compliance team has lived through the pain of assembling audit packages under deadline pressure. The phrase "one-click audit package" triggers an immediate emotional recognition that shortens sales cycles by weeks. **Low technical risk** because most infrastructure (Merkle anchoring, audit log, document storage primitives) already exists in v2.2.

Total estimated effort for Pillar 3: **14 engineer-weeks**. Recommended team: 1 backend engineer + 1 frontend engineer, can be delivered in Q3 alongside Pillar 1 since the two are largely independent. Realistic delivery: Q3 dev → Q4 GA Vault v1.

— 07 • PILLAR 4

Broader Data Coverage

Expands addressable market — each new jurisdiction unlocks 5-15 enterprise prospects in region.

RegGuard AI v2.2 is heavily US/OFAC-centric. Four real feeds are wired in (Federal Register, ESMA, OFAC SDN, EU CFSP), which is enough to demonstrate the concept but caps the addressable market. Enterprise buyers in the UK, Singapore, Hong Kong, Switzerland, Germany, France, and Australia evaluate compliance platforms primarily on whether they cover **their** local regulators. Without that coverage, deals stall at the technical evaluation stage.

4.1 Jurisdiction Expansion: FCA, MAS, SFC, FINMA, BaFin, AMF, ASIC

14 EW

Each new regulator requires identifying the official publication channel (RSS, API, or curated dataset), building a fetcher module compatible with the existing *fetch-free-tier-data.py* architecture, normalizing the data into the existing regwatch schema, and adding jurisdiction-aware filtering to the Regulation Watch view. Estimated 2 engineer-weeks per jurisdiction, batched.

4.2 Sanctions List Expansion: UN, UK HMT, AU DFAT, CA SEMA

6 EW

Extend the existing OFAC + EU CFSP coverage to include the four other major global sanctions lists. UN list is the most strategically valuable because it is the reference list that many other jurisdictions re-publish. All four lists are available as free downloads.

4.3 Enforcement Actions & Court Cases Feed

6 EW

Beyond regulation text, enterprise buyers want to see **how regulators are actually enforcing** the rules. Add a curated feed of enforcement actions (SEC, CFTC, FCA, MAS) and relevant court decisions, normalized into the existing case-management schema.

4.4 Adverse Media Feeds for Entity Screening

4 EW

Negative news screening is a regulatory expectation for KYC/KYB programs. Wire in a curated adverse media feed (GDELT or similar free-tier source) and surface matches alongside sanctions hits in the Sanctions Screening view.

Total estimated effort for Pillar 4: **30 engineer-weeks**. Recommended team: 1 data engineer (primary) + 1 backend engineer (support). Realistic delivery: Q4 this year for first three jurisdictions (FCA, MAS, SFC) → Q1 next year for the remaining four.

— 08 • PILLAR 5

Risk & Screening Automation

Each feature maps to a specific compliance team budget line — sales-led growth opportunity.

This pillar extends RegGuard's surveillance capabilities with features that map directly to budget lines compliance teams already fund. Unlike Pillar 1 (which creates a new category of spend), Pillar 5 features displace existing tools — making budget reallocation the natural sales motion. The Network Graph Explorer v2.2 already provides the foundation for the beneficial ownership graph feature.

5.1 PEP Screening (Politically Exposed Persons)

5 EW

Integrate a curated PEP dataset (free tier from OpenSanctions or similar) and surface matches in the Sanctions Screening view alongside OFAC/CFSP hits. Each match carries a risk tier (e.g. current senior official vs. family member of former official) and a recommended enhanced due diligence workflow.

5.2 Beneficial Ownership Graph (extends Network Graph)

8 EW

Extend the existing force-directed Network Graph Explorer to ingest **corporate ownership chains** from GLEIF LEI data and registry sources. Reveals hidden control relationships — e.g. an apparently clean counterparty whose ultimate beneficial owner appears on a sanctions list three layers up the ownership chain. Builds on the v2.2 enrichment work.

5.3 Anomaly Detection on Transaction Patterns

8 EW

Statistical anomaly detection on transaction streams (unsupervised isolation forest or similar) surfacing structuring patterns, velocity anomalies, and unusual counterparty clusters. Each anomaly carries an explainability score so the human reviewer understands why the model flagged it.

5.4 KYC / KYB Document OCR with Auto-Extracted Fields

7 EW

Upload a passport, utility bill, or incorporation document — RegGuard extracts structured fields (name, DOB, address, registration number) using a free-tier OCR pipeline and pre-populates the KYC/KYB record. Eliminates manual data entry, the most-hated compliance task.

Total estimated effort for Pillar 5: **28 engineer-weeks**. Recommended team: 1 backend engineer + 1 ML engineer. Realistic delivery: Q4 dev → Q1 next year GA for PEP + UBO graph; Q2 next year for anomaly detection and OCR.

— 09 • PILLAR 6

Integration & Public API

Turns RegGuard from product to platform — enables partner ecosystem and 3rd-party integrations.

A public REST API transforms RegGuard from a single product into a platform that other tools integrate with. This unlocks three strategic outcomes: (1) direct revenue via API usage metering, (2) distribution via partner ecosystems (GRC suites, CRMs, SIEMs), and (3) a defensible technical moat. Once enterprises build integrations against the RegGuard API, switching costs lock them in. The API-first pivot is the single highest-leverage strategic move in this roadmap after Pillar 1.

6.1 Public REST API with Usage Metering

8 EW

Versioned REST API exposing regulation search, sanctions screening, risk scoring, and audit anchoring endpoints. Built-in usage metering (per-customer rate limits, monthly call counts) enables usage-based pricing — a monetization model that scales with customer value. Includes developer portal with interactive docs and API keys management.

6.2 Webhook Subscriptions for Real-Time Regulation Alerts

3 EW

Customers register webhook URLs and receive push notifications when regulations relevant to their configured jurisdictions change. Enables downstream automation (auto-create Jira tickets, trigger policy review workflows) without polling.

6.3 GRC Connectors: ServiceNow GRC, Archer, MetricStream

5 EW

Bidirectional connectors to the three major GRC suites. Strategy: position RegGuard as the AI-native intelligence layer that feeds insights into the customer's existing GRC system of record, rather than competing head-on. Co-sell motion with the GRC vendors.

6.4 CRM / ERP Connectors: Salesforce, HubSpot, SAP, Oracle

2 EW

Lightweight connectors (most of the work is API mapping, not deep integration). Enables compliance screening to fire automatically when a new customer or vendor is onboarded in the CRM/ERP — a high-value workflow for sales-led growth.

PLATFORM PLAY

The combination of public API + webhook subscriptions + GRC connectors positions RegGuard as the **AI compliance intelligence layer** that plugs into the customer's existing toolchain. This is the platform positioning that justifies a Series A valuation premium versus a single-product SaaS multiple.

Total estimated effort for Pillar 6: **18 engineer-weeks** (excluding CRM/ERP partner-built connectors). Recommended team: 1 backend engineer + 1 developer-experience engineer. Realistic delivery: Q3 spec → Q4 beta → Q1 next year GA + metering.

— 10 • PILLAR 7

Security & Trust (Enterprise Sales Unlock)

SOC 2 Type II badge alone unlocks 80% of enterprise procurement gates.

Enterprise procurement processes for compliance software routinely require SSO, MFA, and a SOC 2 Type II report before contracts above \$50K ARR can be signed. Without these, deals die at the security review stage regardless of how impressive the product is. This pillar is the single highest-

leverage investment for compressing enterprise sales cycles — every dollar spent here multiplies the close rate of every other pillar's feature.

7.1 SSO (SAML 2.0 + OpenID Connect)

6 EW

Non-negotiable for enterprise buyers. Support both SAML 2.0 (Okta, Azure AD, Google Workspace) and OIDC. Includes just-in-time user provisioning, SCIM for automated user lifecycle management, and tenant-aware role mapping. Without SSO, no enterprise will deploy RegGuard to its compliance team.

7.2 MFA Enforcement + Session Policies

3 EW

Configurable per-tenant MFA requirements (TOTP, WebAuthn), session timeouts, IP allowlists, and concurrent-session limits. Maps directly to the security questionnaire every enterprise buyer sends.

7.3 Data Residency Options (EU / US / APAC)

4 EW

For multi-region enterprises, the ability to pin tenant data to a specific region is a hard requirement. Requires region-aware routing, per-tenant storage selection, and clear data-flow documentation for the security review.

7.4 Customer-Managed Encryption Keys (CMEK)

4 EW

Allow enterprise customers to bring their own encryption keys (AWS KMS, Azure Key Vault, GCP Cloud KMS). Particularly important for financial services customers who must demonstrate key custody to their own regulators.

7.5 SOC 2 Type II Readiness & Audit

3 EW + audit cost

Implement SOC 2 controls (access logging, change management, incident response), engage a CPA firm for the Type II audit, and publish the resulting report under NDA to prospects. The audit itself takes 6-9 months — the work begins in Q3 this year for a Type II report available in Q1 next year.

HIGHEST-LEVERAGE INVESTMENT

SOC 2 Type II badge alone unlocks 80% of enterprise procurement gates. Without it, every enterprise deal requires a custom security review that adds 6-12 weeks to the sales cycle. With it, deals pass through procurement in days. This is the single highest-ROI investment in the entire roadmap for sales cycle compression.

Total estimated effort for Pillar 7: **20 engineer-weeks** plus SOC 2 audit fees (typically 40K—80K). Recommended team: 1 security engineer + external auditor. Realistic delivery: Q3 SSO dev → Q4 SSO GA + SOC 2 audit kickoff → Q1 next year SOC 2 Type II report.

UX Polish

Lower effort, high perceived value — multiplies demo-to-close conversions.

UX polish features are low engineering cost but disproportionately affect perceived product maturity. A demo that includes a command palette, multi-language UI, and drag-and-drop dashboard feels like a mature platform; the same demo without these features feels like a prototype. This pillar is the cheapest way to multiply demo-to-close conversion rates — budget 10 engineer-weeks and ship in Q3 as quick wins alongside the heavier Pillar 1 work.

8.1 Command Palette (Cmd+K)

2 EW

Power-user shortcut to navigate views, search regulations, and run common actions (assign task, export report) without leaving the keyboard. A 2-engineer-week feature that signals product maturity out of proportion to its cost.

8.2 Saved Searches + Custom Alerts

2 EW

Let users save frequent regulation searches and subscribe to email alerts when new results appear. A foundational personalization feature that drives daily-active usage.

8.3 Multi-Language UI (EN / FR / DE / ES / ZH)

2 EW

i18n framework with the four priority languages beyond English. Translation memory per tenant so customers can override standard translations with their internal terminology. Required for European enterprise sales.

8.4 Personalized Dashboard with Drag-and-Drop Widgets

3 EW

Each user composes their own dashboard from a widget library (regulation watchlist, my tasks, risk heat map, recent audit log). Persists per user, not per tenant. Drives daily-active engagement.

8.5 Dark Mode Toggle + Onboarding Wizard

1 EW

Dark mode is already partially supported (this very document is dark themed). Expose the toggle in the user menu. The onboarding wizard reduces time-to-value for new tenants — a 5-step setup flow covering jurisdiction selection, team invite, first policy import.

Total estimated effort for Pillar 8: **10 engineer-weeks**. Recommended team: 1 frontend engineer, sequenced in Q3 as a Quick Win track running parallel to the heavier Pillar 1 work. Realistic delivery: Q3 dev → Q3 GA.

AI Governance & Explainability

Regulator-facing capabilities — uniquely valuable as EU AI Act enforcement begins.

As EU AI Act enforcement phases in over the next 12-18 months, financial services firms using AI for credit scoring, fraud detection, AML screening, or customer segmentation must demonstrate conformity. The governance capabilities in this pillar position RegGuard as the **AI governance ready platform** — uniquely valuable because RegGuard can govern not just its own AI features but also serve as the governance layer for the customer's other AI systems.

The existing “machine proposes, human confirms” design principle is itself a governance feature — it provides a clear audit trail of AI recommendations and the human decisions that confirmed or rejected them. This pillar formalizes and extends that pattern.

9.1 Model Cards Published Per AI Feature

2 EW

Every AI feature ships with a published model card documenting training data, intended use, known limitations, performance metrics across demographic slices, and retraining cadence. Standardized format (Google's model card template is the de-facto industry baseline). Required artifact for AI Act conformity assessments.

9.2 Decision Audit Trail (every recommendation traceable to inputs)

4 EW

For every AI recommendation surfaced in the UI, store the full input context (regulation citations, entity match scores, historical decisions) so the recommendation can be reconstructed and explained months later. Builds on the existing Merkle anchoring infrastructure.

9.3 Human-in-the-Loop Checkpoints (already in v2.2)

0 EW

The *BooleanActionCard* component already implements this pattern across all 29 views — no additional engineering work required. This is a strategic messaging point: RegGuard is already governance-ready where competitors are scrambling to add it.

9.4 Bias Detection in Screening Outcomes

6 EW

Statistical analysis of sanctions screening and risk scoring outcomes across protected attributes (where data is available), surfacing disparate impact patterns. Both a regulatory expectation under AI Act and a moral obligation. Outputs a quarterly bias audit report suitable for board-level review.

AI ACT TIMING

EU AI Act enforcement for high-risk AI systems (which includes financial services AML/fraud screening) phases in over the next 12-18 months. Firms that can demonstrate conformity early will avoid both fines (up to 6% of global revenue) and the operational disruption of retroactive conformity work. **This pillar is timed to that enforcement window.**

Total estimated effort for Pillar 9: **12 engineer-weeks** (excluding the already-shipped human-in-the-loop checkpoint). Recommended team: 1 ML engineer + 1 backend engineer. Realistic delivery: Q3 dev → Q4 GA.


13 • PRIORITIZATION

Impact vs Effort Prioritization Matrix

The matrix below plots all nine enhancement pillars by engineering effort (X-axis, in engineer-weeks) and business impact (Y-axis, on a 1-10 scale). Bubble size indicates revenue potential. The four quadrants — Quick Wins (top-left), Strategic Bets (top-right), Fill-Ins (bottom-left), Defer (bottom-right) — provide the prioritization framework.

IMPACT VS EFFORT MATRIX

Bubble size = revenue potential · Amber = tier-1 priority · Cyan = tier-2/3



Three observations drive the prioritization. **Pillar 7 (Security & SOC 2)** sits in the Quick Wins quadrant with very high impact and moderate effort — ship it first because it unblocks every

enterprise deal. **Pillar 1 (AI Intelligence Layer)** is the largest Strategic Bet — high effort but the deepest moat and biggest revenue uplift. **Pillar 3 (Evidence Vault)** is the sweet spot: low effort, high impact, direct revenue driver — ship it alongside Pillar 7 in Q3.

Pillars 4, 5, and 6 form the second wave — they expand addressable market and enable platform positioning but depend on the Tier 1 foundation being in place. Pillars 2, 8, and 9 are the polish layer — ship them opportunistically when engineering capacity allows.

14 • TIMELINE

Quarterly Delivery Timeline

The Gantt chart below sequences the nine pillars across four quarters (Q3 this year through Q2 next year). Bars are colored by status: amber for active development, cyan for beta, emerald for GA, dashed outline for stretch goals. The sequence prioritizes Tier 1 work in Q3-Q4 to unlock enterprise sales conversations, with Tier 2 and Tier 3 work flowing in parallel where dependencies allow.



Key milestones to plan marketing and sales motions around:

- ▶ **End of Q3:** SSO GA · Evidence Vault v1 · UX polish GA · Public API spec frozen
- ▶ **Mid Q4:** Compliance Copilot v1 GA · Workflow & collaboration GA · SOC 2 audit kickoff
- ▶ **End Q4:** AI governance GA · Public API beta · First new jurisdictions (FCA, MAS, SFC)
- ▶ **Q1 next year:** SOC 2 Type II report published · Public API GA · PEP screening + UBO graph GA
- ▶ **Q2 next year:** Anomaly detection GA · KYC/KYB OCR GA · Remaining jurisdictions (FINMA, BaFin, AMF, ASIC)

— 15 · RECOMMENDATIONS

Top 5 Pre-Market Recommendations

If only five of the nine pillars can be funded before formal market assessment, these are the five. They are ranked by combined impact on (a) AI moat depth, (b) enterprise sales cycle compression, (c) revenue uplift potential, and (d) technical feasibility within the four-quarter window. Each recommendation includes the expected outcome and the investment required.

01

AI Intelligence Layer (Pillar 1)

The deepest moat and biggest revenue uplift. Converts RegGuard from a data dashboard into a reasoning platform with grounded recommendations, gap analysis, and policy drafting. Without this, RegGuard is one of many compliance dashboards; with it, RegGuard is the category-defining AI-native compliance platform. **Invest heaviest here.**

EFFORT 42 EW OUTCOME 2-3x ARR uplift SHIP Q4 Copilot v1 GA

02

Evidence Vault & Audit Packages (Pillar 3)

Direct revenue driver, low technical risk. Most infrastructure already exists in v2.2. "One-click audit package" is the feature that triggers emotional recognition in sales conversations and shortens deal cycles by weeks. Ship alongside Pillar 7 in Q3 as the flagship enterprise sales demo.

EFFORT 14 EW OUTCOME Enterprise deal cycle -40% SHIP Q4 GA Vault v1

03

SSO + SOC 2 Type II Readiness (Pillar 7)

Unblocks enterprise procurement. SOC 2 Type II badge alone opens 80% of enterprise procurement gates; without it, every deal requires a custom security review adding 6-12 weeks to the cycle. Audit takes 6-9 months — start the clock in Q3 this year for a Q1 next year report.

EFFORT 20 EW + audit fees OUTCOME Enterprise close rate +3x SHIP Q1 SOC2 Type II

04

Public REST API + Webhooks (Pillar 6)

Enables platform play. API-first positioning transforms RegGuard from a single product into a platform that other tools integrate with — unlocking partner ecosystem distribution, usage-based monetization, and a defensible technical moat via integration switching costs. Ship beta in Q4, GA in Q1 next year.

EFFORT 18 EW OUTCOME Platform multiple on valuation SHIP Q1 API GA

05

Broader Jurisdiction Coverage (Pillar 4)

Expands addressable market. Each new jurisdiction unlocks 5-15 enterprise prospects in region. Start with FCA (UK), MAS (Singapore), and SFC (Hong Kong) in Q4 — these three cover the largest non-US financial services markets and align with where enterprise buyers already evaluate compliance platforms.

EFFORT 30 EW OUTCOME TAM +60% (UK/SG/HK) SHIP Q4 FCA/MAS/SFC GA

DECISION REQUIRED

Go-ahead requested on Q3 kickoff for Pillars 1, 3, 7, 8 (initial tracks) with parallel Pillar 6 spec work. Total Q3 commitment: approximately 60 engineer-weeks across four parallel tracks. Pillars 2, 4, 5, 9 sequence in Q4 based on Q3 learnings and customer feedback from initial enterprise demos.

The combined investment across all nine pillars is approximately **196 engineer-weeks** of focused development work, distributed across four quarters. With a team of four engineers (two backend, one ML/LLM, one frontend) plus one security engineer added in Q4 for SOC 2 work, this is achievable in the proposed timeline. Total external spend is dominated by the SOC 2 Type II audit fees (\$40-80K) and optional Polygon Amoy testnet MATIC for chain anchoring (free from faucet).

The strategic question for stakeholders is not whether these nine pillars are the right investments — the prioritization matrix and revenue implications make that case clearly. The question is sequencing: which subset can be funded in Q3 to demonstrate momentum ahead of formal market

assessment, with the remaining pillars sequenced based on customer feedback from initial enterprise demos. The recommendation above (Pillars 1, 3, 7, 8 in Q3) is the highest-leverage Q3 combination because it ships the AI moat, the audit-ready evidence vault, the enterprise sales unlock, and the UX polish simultaneously.

CLOSING POSITION

RegGuard AI v2.2 has proven the technical thesis. The nine enhancement pillars in this document convert that proof into an enterprise-grade, investable platform. With Q3 commitment to the recommended Tier 1 tracks, the platform will be ready for formal market assessment by end of Q4 this year, with SOC 2 Type II report arriving in Q1 next year to unlock the full enterprise procurement gate.

From v2.2 Baseline to Investable Platform

RegGuard AI v2.2 has demonstrated technical credibility across 29 state-machine views, four live data feeds, real Merkle anchoring, and a mobile-ready force-directed graph. The nine enhancement pillars in this document convert that technical foundation into an enterprise-grade, investable compliance automation platform. Tier 1 investments — the AI Intelligence Layer, the Evidence Vault, and SOC 2 / SSO readiness — are the highest-leverage moves and should ship in Q3-Q4 this year.